

## ▼ [Aman's AI Journal](#) | Primers | NumPy Tutorial

### Overview

- [NumPy](#) is the core library for scientific computing in Python. It is informally known as the **swiss army knife** of the **data scientist**.
- It provides a high-performance multidimensional array object `numpy.ndarray`, and tools for operating on these arrays.
- If you're already familiar with numerical processing in a different language like MATLAB and R, here are some recommended references:
  - [NumPy for MATLAB users](#)
  - [Python for R users](#)
- **Related primers:** [Matplotlib](#) and [SciPy](#).

### Arrays

- A NumPy array is a **grid of values**, all of the same type, and is indexed by a tuple of non-negative integers.
- The number of dimensions is the **rank** of the array; the **shape** of an array is a tuple of integers giving the size of the array along each dimension.
- We can initialize NumPy arrays from **(nested) lists and tuples**, and access elements using square brackets as array subscripts:

```
import numpy as np
```

```
a = np.array([1, 2, 3])          # Define a rank 1 array using a list
print(type(a))                  # Prints <class 'numpy.ndarray'>
print(a.shape)                  # Prints (3,)
print(a[0], a[1], a[2])         # Prints (1, 2, 3)
a[0] = 5                         # Change an element of the array
print(a)                        # Prints [5 2 3]
```

```
b = np.array([[1, 2, 3]])       # Define a rank 2 array (vector) using a nested list
print(b.shape)                  # Prints (1, 3)
```

B

```

c = np.array([[1, 2, 3], [4, 5, 6]]) # Define a rank 2 array (matrix) using a nested list
print(c.shape)                       # Prints (2, 3)
print(c[0, 0], c[0, 1], c[1, 0])    # Prints (1, 2, 4)

d = np.array((1, 2, 3))              # Define a rank 1 array from a tuple using a nested list
print(d)                             # Prints [1 2 3]
print(d.shape)                       # Prints (3,)

e = np.array(((1, 2, 3), (4, 5, 6))) # Define a rank 2 array using a nested list
print(e)                             # Prints [[1, 2, 3],
                                     #         [4, 5, 6]]

# NumPy arrays can be initialized using other NumPy arrays or lists
# but note that the resulting matrix is always of type NumPy ndarray
l = [1, 2, 3]                        # Define a python list
a = np.array([1, 2, 3])              # Define a numpy array by passing in a list
f = np.array([a, a, a])              # Matrix initialized with NumPy arrays
g = np.array([l, l, l])              # Matrix initialized with lists
h = np.array([a, [1, 1, 1], a])      # Matrix initialized with both types

# All the below statements print [[1 2 3]
#                                 [1 2 3]
#                                 [1 2 3]]
print(f)
print(g)
print(h)

```

- Note the difference between a Python list and a NumPy array. NumPy arrays are designed for numerical (vector/matrix) operations, while lists are for more general purposes.

```
import numpy as np
```

```

l = [1, 2, 3]                # Define a python list
a = np.array([1, 2, 3])      # Define a numpy array by passing in a list
print(l)                    # Prints [1, 2, 3]

```

```
print(a)                # Prints [1 2 3]

print(type(l))          # Prints <class 'list'>
print(type(a))          # Prints <class 'numpy.ndarray'>
```

- Note that when defining an array, be sure that all the rows contain the **same number of columns/elements**. Otherwise, algebraic operations on malformed matrices could lead to unexpected results:

```
import numpy as np

a = np.array([[1, 2], [3, 4]]) # Define a 2x2 matrix

# Print a scaled version of 'a', more on this in the section on "scaling and translating arrays" below
print(a * 2)                  # Prints [[2 4]
                              #          [6 8]]

# Define a malformed matrix. Note the third row contains 3 elements, while other rows contain 2 elements
b = np.array([[1, 2], [3, 4], [5, 6, 7]])

# Print the malformed matrix
print(b)                      # Prints [list([1, 2]) list([3, 4]) list([5, 6, 7])]

# Supposed to scale the whole matrix but does *not*
print(b * 2)                  # Prints [list([1, 2, 1, 2]) list([3, 4, 3, 4]) list([5, 6, 7, 5, 6, 7])]
```

- NumPy also provides many functions to create arrays:

```
import numpy as np

a = np.zeros((2, 2))          # Define an array of all zeros
print(a)                      # Prints [[ 0.  0.]
                              #          [ 0.  0.]
```

```

b = np.ones((1, 2))          # Define an array of all ones
print(b)                     # Prints [[ 1.  1.]]

c = np.full((2, 2), 7)       # Define a constant array
print(c)                     # Prints [[ 7.  7.]
                             #          [ 7.  7.]]

d = np.eye(2)                # Define a 2x2 identity matrix
print(d)                     # Prints [[ 1.  0.]
                             #          [ 0.  1.]]

e = np.random.random((2, 2)) # Define a 2x2 matrix from the uniform distribution [0, 1)
print(e)                     # Prints a 2x2 matrix of random values

g = 5 * np.random.random_sample((2, 2)) - 5 # Sample 2x2 matrix from Unif[-5, 0)
print(f)                     # Sample from Unif[a, b), b > a: (b - a) * random_sample() + a
                             # Prints a 2x2 matrix of random values

f = np.random.randn(2, 2)    # Sample a 2x2 matrix from the "standard normal" distribution
print(f)                     # Prints a 2x2 matrix of random values

g = 2.5 * np.random.randn(2, 2) + 3 # Sample 2x2 matrix from N(mean=3, var=6.25)
print(f)                     # General form: stddev * np.random.randn(...) + mean
                             # Prints a 2x2 matrix of random values

```

- Note that with `np.random.randn()`, the length of each dimension of the output array is an **individual argument**. On the other hand, `np.random.random()` accepts its shape argument as a **single tuple containing all dimensions**. More on this in the section on [standard normal](#).
- You can read about other methods of array creation in the [NumPy documentation](#).

## ▼ Array indexing

- NumPy offers several ways to index into arrays.

### Row and column indexing

- To "select" a particular row or column in an array, NumPy offers similar functionality as Python lists:

```
import numpy as np

a = np.array([[1, 2], [3, 4]])

# Select a row
a[0]          # Prints [1, 2]
a[1]          # Prints [3, 4]

# Select a column
a[:, 0]       # Prints [1, 3]
a[:, 1]       # Prints [2, 4]
```

- You can use `np.arange()` to select the rows/columns of an array. For more details on `np.arange()`, refer to the section on [arange](#) below.

```
import numpy as np

a = np.array([[1, 2], [3, 4]])

# Return the entire array
a[np.arange(2), :] # Prints [[1, 2],
                    #         [3, 4]]

# Return the first row
a[np.arange(1), :] # Prints [[1, 2]]
```

- You can use an **"index-array"** that contains indices of rows or columns to index into another array. This is a very **common use-case** in NumPy-based projects.

```
import numpy as np
```

```

a = np.array([[1, 2], [3, 4]])

# Selecting columns using an index array
b = [0, 0]          # Select the first column for both rows (see below)
a[np.arange(1), b] # Prints [1, 1] (same as a[0, b])
a[np.arange(2), b] # Prints [1, 3] (same as a[[0, 1], b])

a[:, b]            # Prints [[1, 1],
                    #         [3, 3]]

# Selecting rows using an index array
b = [0, 0]          # Select the first row for both columns (see below)
a[b, np.arange(1)] # Prints [1, 1] (same as a[b, 0])
a[b, np.arange(2)] # Prints [1, 2] (same as a[b, [0, 1]])

a[b, :]            # Prints [[1, 2],
                    #         [1, 2]]

```

## ▼ Slicing

- Similar to Python lists, NumPy arrays can be sliced.
- Since arrays may be multidimensional, you must specify a slice for each dimension of the array:

```

import numpy as np

# Define the following rank 2 array with shape (3, 4)
# [[ 1  2  3  4]
#  [ 5  6  7  8]
#  [ 9 10 11 12]]
a = np.array([[1, 2, 3, 4], [5, 6, 7, 8], [9, 10, 11, 12]])

# Use slicing to pull out the subarray consisting of the first 2 rows
# and columns 1 and 2; b is the following array of shape (2, 2):
# [[2 3]

```

```
# [6 7]]
b = a[:2, 1:3]

# A slice of an array is a view into the same data, so modifying it
# will modify the original array
print(a[0, 1]) # Prints 2
b[0, 0] = 77   # b[0, 0] is the same piece of data as a[0, 1]
print(a[0, 1]) # Prints 77
```

- You can also mix integer indexing with slice indexing. However, doing so will yield an array of lower rank than the original array.
- Note that this is quite different from the way that MATLAB handles array slicing:

```
import numpy as np

# Define the following rank 2 array with shape (3, 4)
# [[ 1  2  3  4]
#  [ 5  6  7  8]
#  [ 9 10 11 12]]
a = np.array([[1, 2, 3, 4], [5, 6, 7, 8], [9, 10, 11, 12]])

# Two ways of accessing the data in the middle row of the array.
# Mixing integer indexing with slices yields an array of lower rank,
# while using only slices yields an array of the same rank as the
# original array:
row_r1 = a[1, :]           # Rank 1 view of the second row of a
row_r2 = a[1:2, :]         # Rank 2 view of the second row of a
print(row_r1, row_r1.shape) # Prints [5 6 7 8] (4,)
print(row_r2, row_r2.shape) # Prints [[5 6 7 8]] (1, 4)

# We can make the same distinction when accessing columns of an array:
col_r1 = a[:, 1]
col_r2 = a[:, 1:2]
print(col_r1, col_r1.shape) # Prints [ 2  6 10] (3,)
print(col_r2, col_r2.shape) # Prints [[ 2]
#                               [ 6]
#                               [10]] (3, 1)
```

```
# Mix row and column slicing to print the first 2 rows and alternate
# columns
arr_row_col = a[:2, ::2]
print(arr_row_col, arr_row_col.shape) # Prints [[1, 3],
                                     #          [5, 7]] (2, 2)
```

## ▼ Integer array indexing

- When you index into NumPy arrays using slicing, the resulting array view will always be a subarray of the original array.
- In contrast, integer array indexing allows you to construct arbitrary arrays using the data from another array.
- Here is an example:

```
import numpy as np

a = np.array([[1, 2], [3, 4], [5, 6]])

# An example of integer array indexing.
# The returned array will have shape (3,) and
print(a[[0, 1, 2], [0, 1, 0]]) # Prints [1 4 5]

# The above example of integer array indexing is equivalent to this:
print(np.array([a[0, 0], a[1, 1], a[2, 0]])) # Prints [1 4 5]

# When using integer array indexing, you can reuse the same
# element from the source array:
print(a[[0, 0], [1, 1]]) # Prints [2 2]

# Equivalent to the previous integer array indexing example
print(np.array([a[0, 1], a[0, 1]])) # Prints [2 2]
```

- One useful trick with integer array indexing is selecting or mutating one element from each row of a matrix:



```
import numpy as np

# Define a new array from which we will select elements
a = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])

print(a)                # Prints [[ 1,  2,  3],
                        #           [ 4,  5,  6],
                        #           [ 7,  8,  9],
                        #           [10, 11, 12]]

# Define an array of indices
b = np.array([0, 2, 0, 1])

# Select one element from each row of a using the indices in b
print(a[np.arange(4), b]) # Prints [ 1  6  7 11]

# Mutate one element from each row of a using the indices in b
a[np.arange(4), b] += 10

print(a)                # Prints [[11,  2,  3],
                        #           [ 4,  5, 16],
                        #           [17,  8,  9],
                        #           [10, 21, 12]]
```

## ▼ Boolean array indexing

- Boolean array indexing lets you pick out arbitrary elements of an array.
- Frequently this type of indexing is used to select the elements of an array that satisfy some condition.
- Here is an example:

```
import numpy as np

a = np.array([[1, 2], [3, 4], [5, 6]])
```

```

print(a[True])      # Same as print(a), interpreted as a "True" mask
                   # on each of a's elements

bool_idx = (a > 2) # Find the elements of a that are bigger than 2;
                 # this returns a NumPy array of Booleans of the same
                 # shape as a, where each slot of bool_idx tells
                 # whether that element of a is > 2

print(bool_idx)     # Prints [[False False]
                   #          [ True  True]
                   #          [ True  True]]

# We use boolean array indexing to construct a rank 1 array
# consisting of the elements of a corresponding to the True values
# of bool_idx
print(a[bool_idx]) # Prints [3 4 5 6]

# We can do all of the above in a single concise statement:
print(a[a > 2])     # Prints [3 4 5 6]

```

- As an extension of the above concept, to select elements on an array based on the elements of another array:

```

import numpy as np

a = np.array([1, 2, 3, 4, 5, 6])
b = np.array(['f','o','o','b','a','r'])

# Note that & is the bitwise AND operator, and && is not supported in NumPy,
# so np.logical_and() is used in the below example
print(b[np.logical_and((a > 1), (a < 5))]) # Prints ['o','o','b']

# Another way to accomplish this is using np.all(),
# which is explained in the section on "all" below.
print(b[np.all([a > 1, a < 5], axis=0)])   # Prints ['o','o','b']

```

- For brevity, we have left out a lot of details about NumPy array indexing; if you want to know more you should read the [NumPy documentation](#).

## ► Datatypes

- Every NumPy array is a grid of elements of the same type.
- NumPy provides a large set of numeric datatypes that you can use to construct arrays.
- NumPy tries to guess a datatype when you create an array, but functions that construct arrays usually also include an optional argument to explicitly specify the datatype.
- As an example:

[ ] ↳ 1 cell hidden

- You can read all about datatypes in the [NumPy documentation](#).

## ▼ Array math

- NumPy provides many functions for manipulating arrays; you can find an exhaustive list in the [NumPy documentation](#). Some common ones are listed below.

### Element-wise operations

- Algebraic operations such as  $+$ ,  $-$ ,  $*$ ,  $/$  etc. are available both as operator overloads and as functions in the NumPy module. These operators carry out **element-wise** operations on NumPy arrays:

```
import numpy as np
```

```
x = np.array([[1, 2], [3, 4]], dtype=np.float64)
```

```
y = np.array([[5, 6], [7, 8]], dtype=np.float64)
```

```
# Elementwise sum; both produce the array
```

B

```
# [[ 6.0  8.0]
# [10.0 12.0]]
print(x + y)
print(np.add(x, y))

# Elementwise difference; both produce the array
# [[-4.0 -4.0]
# [-4.0 -4.0]]
print(x - y)
print(np.subtract(x, y))

# Elementwise product; both produce the array
# [[ 5.0 12.0]
# [21.0 32.0]]
print(x * y)
print(np.multiply(x, y))

# Elementwise division; both produce the array
# [[ 0.2      0.33333333]
# [ 0.42857143  0.5      ]]
print(x / y)
print(np.divide(x, y))

# Elementwise square root; produces the array
# [[ 1.      1.41421356]
# [ 1.73205081  2.      ]]
print(np.sqrt(x))
```

## ▼ NumPy arrays vs. Python lists

- A common mistake is to mix up the concepts of NumPy arrays and Python lists.
- Note that the `+` operator in the context of NumPy arrays performs an **element-wise addition**, while the same operation on Python lists results in a **list extension**.

```
import numpy as np

l = [1, 2, 3]          # Define a python list
a = np.array([1, 2, 3]) # Define a numpy array by passing in a list
print(a + a)           # Prints [2 4 6]
print(l + l)           # Prints [1, 2, 3, 1, 2, 3]
```

- With NumPy arrays, we can **scale** the vector with `*` by performing **element-wise multiplication**, while the same operation on Python lists results in a **list concatenation** (and in MATLAB, results in **matrix multiplication**).
  - We instead use the `dot` function to compute inner products of vectors, to multiply a vector by a matrix, and to multiply matrices. We discuss this in detail in the section on [dot product](#) below.

```
import numpy as np

l = [1, 2, 3]          # Define a python list
a = np.array([1, 2, 3]) # Define a numpy array by passing in a list
print(a * 3)           # Prints [3 6 9]
print(l * 3)           # Prints [1, 2, 3, 1, 2, 3, 1, 2, 3]
```

- Make note of this while coding since a function can utilize both Python lists and NumPy arrays. Knowing this can save many headaches!

## ▼ Scaling and translating arrays

- Using regular algebraic operators like `+` and `-` that we saw in the prior section on [element-wise operations](#), we can scale and translate/shift NumPy arrays.
- Operations can be performed between:
  - Two NumPy arrays or,
  - NumPy arrays and scalars.

```
import numpy as np
```

```
a = np.array([[1, 2], [3, 4]]) # Define a 2x2 matrix

# Scale by 2 units and translate by 1 unit
result = a * 2 + 1           # Multiply each element in the matrix by 2 and add 1
print(result)                # Prints [[3 5]
                             #          [7 9]]
```

## ▼ Norm

- NumPy offers a set of algebraic functions in the class **linalg**, which includes the **norm** function:

```
import numpy as np

a = np.array([1, 2, 3, 4])    # Define an array
norm1 = np.linalg.norm(a)

a = np.array([[1, 2], [3, 4]]) # Define a 2x2 matrix
norm2 = np.linalg.norm(a)

# Both print 5.477225575051661
print(norm1)
print(norm2)
```

- By default, `np.linalg.norm()` calculates the **Frobenius norm** for matrices and the **2-norm** for vectors, unless the `ord` argument is overridden.
  - Recall from linear algebra that the **norm** (or magnitude) of an n-dimensional vector  $v$  is the **square root of the sum of its elements squared**:

```
$$
norm(\vec{v}) = ||\vec{v}||_2 = \sqrt{\sum_{i=1}^n v_i^2}
$$
```

- Also, the **Frobenius norm** is the generalization of the vector norm for matrices, and is defined for a matrix  $A$  as:

$$\|\mathrm{A}\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}$$

- The following code snippet compares a manual implementation of the norm vs. the `np.linalg.norm()` function. Both yield the same result.

```
import numpy as np

a = np.array([[1, 2], [3, 4]]) # Define a 2x2 matrix

norm = np.sqrt(np.sum(np.square(a)))

# Both should yield the same result and print:
# 5.477225575051661 == 5.477225575051661
print(norm, '== ', np.linalg.norm(a))

a = np.array([1, 2, 3, 4]) # Define an array

norm = np.sqrt(np.sum(np.square(a)))

# Both should yield the same result and print:
# 5.477225575051661 == 5.477225575051661
print(norm, '== ', np.linalg.norm(a))
```

- Note that if the `axis` argument to `np.linalg.norm()` is not explicitly specified, the function defaults to treating the matrix as a "flat" array of numbers implying that it takes every element of the matrix into account. However, it is possible to get the a row-wise or column-wise norm using the `axis` parameter:
  - `axis=0` operates across all rows, i.e., gets the norm of each **column**.
  - `axis=1` operates across all columns, i.e., gets the norm of each **row**.

```
import numpy as np

a = np.array([[1, 1], [2, 2], [3, 3]]) # Define a 3x2 matrix

normByCols = np.linalg.norm(a, axis=0) # Get the norm for each column; returns 2 elements
normByRows = np.linalg.norm(a, axis=1) # get the norm for each row; returns 3 elements

print(normByCols)          # [3.74165739 3.74165739]
print(normByRows)          # [1.41421356 2.82842712 4.24264069]
```

## ▼ Dot product

- Dot product is available both as a function in the NumPy module `np.dot()` and as an instance method of array objects `<ndarray>.dot()`:

```
import numpy as np

x = np.array([[1, 2], [3, 4]])
y = np.array([[5, 6], [7, 8]])

v = np.array([9, 10])
w = np.array([11, 12])

# Dot/scalar/inner product of vectors; both produce a scalar: 219
print(v.dot(w))
print(np.dot(v, w))

# Matrix/vector product; both produce the rank 1 array [29 67]
print(x.dot(v))
print(np.dot(x, v))

# Matrix/matrix product; both produce the rank 2 array
# [[19 22]
#  [43 50]]
print(x.dot(y))
print(np.dot(x, y))
```



- Some alternative ways to obtain the dot product:

```
import numpy as np

print(np.sum(v * w)) # Using the definition of dot product

print(v @ w)          # numpy>=1.10 overloads the new Pythonic operator "@" for dot product

# Inefficient, unrolled, non-vectorized implementation
dotProduct = 0
for a, b in zip(v, w):
    dotProduct += a * b

print(dotProduct)
```

- `np.dot()` is strongly recommended since it accepts both NumPy arrays and Python lists:

```
import numpy as np

n1 = np.dot(np.array([1, 2]), np.array([3, 4])) # Dot product on numpy arrays
n2 = np.dot([1, 2], [3, 4])                    # Dot product on python lists

# Both print 11
print(n1)
print(n2)
```

## ▼ Broadcasting

- Broadcasting is a powerful mechanism that allows NumPy to work with arrays of different shapes when performing arithmetic operations. Frequently, we have a smaller array and a larger array, and we want to use the smaller array multiple times to perform some operation on the larger array.

B

- Broadcasting typically makes your code more concise and faster, so you should strive to use it where possible.
- The simplest broadcasting example occurs when an array and a scalar value are combined in an operation:

```
import numpy as np

a = np.array([1, 2, 3])
b = 2
print(a * b) # Prints [ 2, 4, 6]

# Note that the result is equivalent to the next example where b is an array!
a = np.array([1, 2, 3])
b = np.array([2, 2, 2])
print(a * b) # Prints [ 2, 4, 6]
```

- Consider another example where we'd like to add a constant vector to each row of a matrix. We could do it like this:

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])
v = np.array([1, 0, 1])
y = np.empty_like(x) # Define an empty matrix with the same shape as x

# Add the vector v to each row of the matrix x with an explicit loop
for i in range(4):
    y[i, :] = x[i, :] + v

# Now y is the following
# [[ 2  2  4]
#  [ 5  5  7]
#  [ 8  8 10]
#  [11 11 13]]
print(y)
```

- This works; however when the matrix  $x$  is very large, computing an explicit loop in Python could be slow.
- Note that adding the vector  $v$  to each row of the matrix  $x$  is equivalent to forming a matrix  $vv$  by stacking multiple copies of  $v$  vertically, then performing element-wise summation of  $x$  and  $vv$ . We could implement this approach like this:

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])
v = np.array([1, 0, 1])
vv = np.tile(v, (4, 1)) # Stack 4 copies of v on top of each other
print(vv)               # Prints [[1 0 1]
                        #           [1 0 1]
                        #           [1 0 1]
                        #           [1 0 1]]
y = x + vv               # Add x and vv element-wise
print(y)                 # Prints [[ 2  2  4
                        #           [ 5  5  7]
                        #           [ 8  8 10]
                        #           [11 11 13]]
```

- NumPy broadcasting allows us to perform this computation without actually creating multiple copies of  $v$ . Consider this version, using broadcasting:

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])
v = np.array([1, 0, 1])
y = x + v # Add v to each row of x using broadcasting
print(y)  # Prints [[ 2  2  4]
          #           [ 5  5  7]
```

```
#      [ 8  8 10]
#      [11 11 13]]
```

- The line  $y = x + v$  works even though  $x$  has shape  $(4, 3)$  and  $v$  has shape  $(3,)$  due to broadcasting; this line works as if  $v$  actually had shape  $(4, 3)$ , where each row was a copy of  $v$ , and the sum was performed element-wise.

## ▼ Rules

- Broadcasting two arrays together follows these rules:
  - If the arrays **do not** have the **same rank**, prepend the shape of the lower rank array with ones until both shapes have the same length.
    - This implies that the arrays do not need to have the same number of dimensions.
  - If the arrays have the **same rank**, NumPy compares their shapes element-wise. It starts with the **trailing dimensions** and works its way forward, checking if the arrays are **compatible** in every dimension, as follows:
    - The two arrays are said to be compatible in a dimension if:
      - They have the **same size** in the dimension or,
      - One of the arrays has **size 1** in that dimension.
    - The arrays can only be broadcast together if they are compatible in all dimensions.
    - When either of the dimensions belonging to both arrays being compared is 1, the **other** is used. In other words, dimensions with size 1 are **stretched or “copied”** to match the other. This is the **primary broadcasting** step.
    - After broadcasting, each array behaves as if its shape is now equal to the element-wise maximum of shapes of the two input arrays.
- Functions that support broadcasting are known as **universal functions**. You can find an exhaustive list of universal functions in the [NumPy documentation](https://numpy.org/doc/stable/user/basics.broadcasting.html).

## Examples

- Examples from [NumPy's documentation on Broadcasting](#):

1. **\*\*Either array has axes with length 1\*\***: In the following example, array `\\(B\\)` has an axis with length `\\(1\\)`.
  - To tackle an axis with length `\\(1\\)` during broadcasting, we simply expand that dimension to match the other array's.

```

...
A      (3d array):  15 x 3 x 5
B      (3d array):  15 x 1 x 5
Result (3d array):  15 x 3 x 5
...

```

1. **\*\*Both arrays have axes with length 1\*\***: In the following example, arrays `\\(A\\)` and `\\(B\\)` both have axes with length `\\(1\\)`.
  - Extrapolating the case above, to tackle an axis with length `\\(1\\)` in either array during broadcasting, we simply expand that

```

...
A      (4d array):  8 x 1 x 6 x 1
B      (3d array):   7 x 1 x 5
Result (4d array):  8 x 7 x 6 x 5
...

```

1. **\*\*Rank mismatch\*\***: In the following example, arrays `\\(A\\)` and `\\(B\\)` do not have the same rank.
  - Recall that to tackle rank mismatch during broadcasting, we prepend the shape of the lower rank array with ones until both shapes

```

...
A      (2d array):  5 x 4
B      (1d array):   4
Result (2d array):  5 x 4
...

```

```

...
A      (3d array):  15 x 3 x 5
B      (2d array):   3 x 5
Result (3d array):  15 x 3 x 5
...

```

```

'''

```

```

A      (3d array): 256 x 256 x 3
B      (1d array):          3
Result (3d array): 256 x 256 x 3
'''

```

1. **\*\*Rank mismatch and axes with length 1\*\***: In the following example, arrays `A` and `B` do not have the same rank and `B` - Tackling rank mismatch was discussed in example `(3)`. Also, handling an axis with length `(1)` was discussed in example `(1)`

```

'''

```

```

A      (2d array):  5 x 4
B      (1d array):    1
Result (2d array):  5 x 4
'''

```

```

'''

```

```

A      (3d array): 15 x 3 x 5
B      (2d array):   3 x 1
Result (3d array): 15 x 3 x 5
'''

```

- Some examples that **do not broadcast**:

1. **\*\*Dimension mismatch\*\***: In the following example, arrays `A` and `B` are not compatible in each dimension and thus, cannot - Recall that the two arrays need to be **\*\*compatible\*\*** in each dimension (starting from the trailing dimension) for broadcasting

```

'''

```

```

A      (1d array):  3
B      (1d array):  4          # Trailing dimensions do not match
'''

```

```

'''

```

```

A      (2d array):    2 x 1
B      (3d array):  8 x 4 x 3 # Second from last dimensions mismatched
'''

```

- Broadcasting examples in code:

```
import numpy as np

# Compute outer product of vectors
v = np.array([1, 2, 3]) # v has shape (3,)
w = np.array([4, 5])    # w has shape (2,)
# To compute an outer product, we first reshape v to be a column
# vector of shape (3, 1); we can then broadcast it against w to yield
# an output of shape (3, 2), which is the outer product of v and w:
# [[ 4  5]
#  [ 8 10]
#  [12 15]]
print(np.reshape(v, (3, 1)) * w)
# Note that np.reshape() is described in detail in the section on "reshape" below.

# Add a vector to each row of a matrix
x = np.array([[1, 2, 3], [4, 5, 6]])
# x has shape (2, 3) and v has shape (3,) so they broadcast to (2, 3),
# giving the following matrix:
# [[2 4 6]
#  [5 7 9]]
print(x + v)

# Add a vector to each column of a matrix
# x has shape (2, 3) and w has shape (2,)
# If we transpose x then it has shape (3, 2) and can be broadcast
# against w to yield a result of shape (3, 2); transposing this result
# yields the final result of shape (2, 3) which is the matrix x with
# the vector w added to each column. Gives the following matrix:
# [[ 5  6  7]
#  [ 9 10 11]]
print((x.T + w).T)
# Another solution is to reshape w to be a column vector of shape (2, 1);
# we can then broadcast it directly against x to produce the same
```

```
# output.
print(x + np.reshape(w, (2, 1)))

# Multiply a matrix by a constant:
# x has shape (2, 3). NumPy treats scalars as arrays of shape ();
# these can be broadcast together to shape (2, 3), producing the
# following array:
# [[ 2  4  6]
#  [ 8 10 12]]
print(x * 2)
```

## ▼ Data centering

- In some scenarios, centering the elements of a dataset is an important preprocessing step. Centering a dataset involves **subtracting the mean from the data**. The resultant data is called zero-centered, since the mean of the resultant dataset is 0.

### Column-centering

- Column-centering a matrix involves subtracting the column mean from each element within a particular column. Note that the sum by columns of a centered matrix is always 0.

```
import numpy as np

a = np.array([[1, 1], [2, 2], [3, 3]]) # Define a 3x2 matrix.

centered = a - np.mean(a, axis=0)      # Remove the mean for each column

# Column-centered matrix
print(centered)                        # Prints [[-1. -1.]
                                     #          [ 0.  0.]
                                     #          [ 1.  1.]]

# New mean-by-column should be 0
print(centered.mean(axis=0))           # Prints [0.  0.]
```

B



## ▼ Row-centering

- For row centering, transpose the matrix, center by columns, and then transpose back the result.

```
import numpy as np

a = np.array([[1, 3], [2, 4], [3, 5]]) # Define a 3x2 matrix.

centered = a.T - np.mean(a, axis=1)    # Remove the mean for each row
centered = centered.T                  # Transpose back the result

# Row-centered matrix
print(centered)                        # Prints [[-1.  1.]
                                         #         [-1.  1.]
                                         #         [-1.  1.]]

# New mean-by-rows should be 0
print(centered.mean(axis=1))           # Prints [0.  0.  0.]
```

## ▼ Selected functions

- NumPy provides a host of useful functions for performing computations on arrays. Below, we've touched upon some of the most useful ones that you'll encounter regularly in projects.
- You can find an exhaustive list of mathematical functions in the [NumPy documentation](#).
- As discussed in the section on [norm](#), most of these functions accept an `axis` parameter which when specified, performs the operation on a per-column (`axis=0`) or per-row basis (`axis=1`) rather than on the entire array.

## Initialization

### Random seed

- Provide a seed to the generator. Used to foster reproducibility of random numbers.

B

```
import numpy as np

np.random.seed(0) # Here, 0 is the seed.

a = np.random.rand()
print(a)          # Print 0.5488135039273248

np.random.seed(0) # Set seed again for next random number to match the prior one.

a = np.random.rand()
print(a)          # Print 0.5488135039273248
```

## ▼ Standard normal

- Return samples from the standard normal (also called the standard Gaussian) distribution.
- Note that `np.random.randn()` is a convenience function for users porting code from MATLAB, and wraps `np.random.standard_normal()`. As such, `np.random.randn()` takes **dimensions as individual ints**.
  - On the other hand, `np.random.standard_normal()` takes a `shape` argument as a **tuple** to specify the size of the output, which is consistent with other NumPy functions like `np.zeros()` and `np.ones()`.

```
import numpy as np

a = np.random.randn(2, 2)          # Sample a 2x2 matrix from the standard normal distribution
print(a)                          # Prints a 2x2 matrix of random values

b = np.random.standard_normal((2, 2)) # Sample a 2x2 matrix from the standard normal distribution
print(b)                          # Prints a 2x2 matrix of random values
```

- To sample the normal distribution  $a \sim \mathcal{N}(\mu, \sigma^2)$ , multiply the output of `np.random.randn()` by  $\sigma$  and add  $\mu$ , i.e., `stddev * np.random.randn(...) + mean`

```
import numpy as np

a = 2.5 * np.random.randn(2, 2) + 3 # Sample 2x2 matrix from N(mean=3, var=6.25)
                                     # General form: stddev * np.random.randn(...) + mean
print(a)                           # Prints a 2x2 matrix of random values
```

## ▼ Uniform

- Samples random floats from a continuous uniform distribution – the half-open interval  $[0.0, 1.0)$ .
- Note that `np.random.rand()` is a convenience function for users porting code from MATLAB, and wraps `np.random.random()`. As such, `np.random.rand()` takes **dimensions as individual ints**.
  - On the other hand, `np.random.random()` takes a `shape` argument as a **tuple** to specify the size of the output, which is consistent with other NumPy functions like `np.zeros()` and `np.ones()`.

```
import numpy as np

a = np.random.rand(2, 2)      # Define a 2x2 matrix from the uniform distribution [0, 1)
print(a)                      # Prints a 2x2 matrix of random values

b = np.random.random((2, 2)) # Define a 2x2 matrix from the uniform distribution [0, 1)
print(b)                      # Prints a 2x2 matrix of random values
```

- To sample the uniform distribution  $a \sim \text{Unif}[x, y)$ , given  $y > x$ , Note that `np.random.uniform()` offers this functionality (drawing samples from a uniform distribution within an arbitrary range) much more directly by accepting three arguments: `low=0.0`, `high=1.0`, `size=None`. The following code snippet thus yields the same output as the one above.

```
import numpy as np

a = np.random.uniform(-5, 0, size=(2,2)) # Sample 2x2 matrix from Unif[-5, 0)

print(a)                                # Prints a 2x2 matrix of random values
```

- Note that `np.random.random()` also offers this functionality (drawing samples from a uniform distribution within an arbitrary range) by multiplying its output by  $(y - x)$  and adding  $x$ , i.e.,  $(y - x) * np.random.random(...) + x$

```
import numpy as np

a = 5 * np.random.random((2, 2)) - 5 # Sample 2x2 matrix from Unif[-5, 0)
                                     # Sample from Unif[x, y), y > x: (y - x) * np.random.random_sample() + x
print(a)                             # Prints a 2x2 matrix of random values
```

- Without explicit arguments, the functions `rand()`, `random()`, `uniform()` and `random_sample()` are equivalent, producing a random float in the range  $[0.0, 1.0)$ .

```
import numpy as np

np.random.seed(0)

a = np.random.rand()
print(a)                # Print 0.5488135039273248

b = np.random.random()
print(b)                # Print 0.7151893663724195

c = np.random.uniform()
print(c)                # Print 0.6027633760716439

d = np.random.random_sample() # numpy.random.random() is an alias for numpy.random.random_sample()
print(d)                # Print 0.5448831829968969
```

## ▼ Random choice

- Generates a random sample from a given 1D array. The main arguments to `np.random.choice()` are `a` and `size`.

B

- `a` can be an `ndarray`, in which case a random sample is returned from its elements. If it's an int, the random sample is generated as if `a` were `np.arange(a)`.
- `size` is an optional argument that can hold an int or a tuple of ints. If it is not overridden, a single value is returned by default.

```
import numpy as np

# Generate a uniform random sample from np.arange(5) of size 3
# This is equivalent to np.random.randint(0,5,3)
print(np.random.choice(5, 3))           # Prints [0 3 4]

# Generate a non-uniform random sample from np.arange(5) of size 3
print(np.random.choice(5, 3, p=[0.1, 0, 0.3, 0.6, 0])) # Prints [3 3 0]

# Generate a uniform random sample from np.arange(5) of size 3 without replacement
print(np.random.choice(5, 3, replace=False))          # Prints [4 1 3]
```

## ▼ Arange

- Return evenly spaced values within the half-open interval  $[start, stop)$  (in other words, the interval including start but excluding stop).
- For integer arguments the function is equivalent to the Python built-in range function, but returns an `ndarray` rather than a list.

```
import numpy as np

print(np.arange(8))           # Prints [0 1 2 3 4 5 6 7]
print(np.arange(3, 8))       # Prints [3 4 5 6 7]
print(np.arange(3, 8, 2))    # Prints [3 5 7]
```

- When using a non-integer step, such as 0.1, the results will often not be consistent. It is better to use `np.linspace()` for those cases as below.

## ▼ Linspace

B

- Return evenly spaced numbers over a specified interval.
- Returns 50 evenly spaced samples (which can be overridden by `num`), calculated over the interval  $[start, stop]$ .

```
import numpy as np

print(np.linspace(2.0, 3.0, num=5))          # Prints [2.   2.25 2.5   2.75 3.   ]
print(np.linspace(2.0, 3.0, num=5, endpoint=False)) # Prints [2.   2.2  2.4  2.6  2.8]
```

## ▼ Reshape

- Gives a new shape to an array without changing its data.
- Note that `np.reshape()` returns a **view** of the array when the elements are contiguous in memory (just like [np.ravel\(\)](#)). So modifying the result of `np.reshape()` would also modify the original array. However, `np.reshape()` returns a copy if, for e.g., the input array were made from slicing another array using a **non-unit** step size (e.g. `a = x[:, :2]`).

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])

print(np.reshape(a, (3, 2)) # Prints [[1 2]
                                #       [3 4]
                                #       [5 6]]

print(np.reshape(a, 6))      # Prints [1  2  3  4  5  6]
```

## ▼ -1 in Reshape

- NumPy allows us to assign **one** (and only one) value of the shape argument to a function as `-1`. NumPy figures out the unknown dimension by looking at the length of the array and the remaining dimensions, while making sure it satisfies the criterion that **the new shape should be compatible with the original shape**.

B

```

import numpy as np

a = np.array([[1, 2, 3, 4], [5, 6, 7, 8], [9, 10, 11, 12]])
print(a.shape)  # Prints (3, 4)

# Reshaping with *just* (-1) leads to array flattening.
# Yields new shape as (12,), a rank 1 array which is compatible with original shape (3, 4)
print(a.reshape(-1))  # Prints [ 1  2  3  4  5  6  7  8  9 10 11 12]
# Note that np.reshape(-1) returns the same result as np.flatten()
# with the exception that np.flatten() returns a copy while reshape() returns a view of the original array.
# For more on np.flatten(), refer the section on "flatten" below.

# Transform a matrix to a row vector by reshaping with (1, -1), i.e., 1 row and 3*4 number of columns.
# Yields new shape as (1, 12), a rank 2 array which is compatible with original shape (3, 4)
print(a.reshape(1, -1)) # Prints [[ 1  2  3  4  5  6  7  8  9 10 11 12]]

# Transform a matrix to a column vector by reshaping with (-1, 1), i.e., 3*4 number of rows and 1 column.
# Yields new shape as (12, 1), a rank 2 array which is compatible with original shape (3, 4)
print(a.reshape(-1, 1)) # Prints [[ 1],
#                        [ 2],
#                        [ 3],
#                        [ 4],
#                        [ 5],
#                        [ 6],
#                        [ 7],
#                        [ 8],
#                        [ 9],
#                        [10],
#                        [11],
#                        [12]]

```

## ▼ Flatten

- Return a **copy** of the array collapsed into one dimension.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])
print(a.flatten()) # Prints [1, 2, 3, 4, 5, 6]
```

## ▼ Squeeze

- Remove single-dimensional entries from the shape of an array.

```
import numpy as np

a = np.array([[[0], [1], [2]]])

print(a.shape)          # Prints (1, 3, 1)
print(np.squeeze(a))     # Prints [0 1 2]
print(np.squeeze(a).shape) # Prints (3,)
```

## ▼ Copy

- Return an array copy of the given object.

```
import numpy as np

a = np.array([[1, 2], [3, 4]])
b = a
c = np.copy(a)

a[0] = 10

print(a)      # Prints [[10 10]
              #          [ 3  4]]
```



```
a[0] == b[0] # Prints True
a[0] == c[0] # Prints False
```

- Note that `np.copy()` performs a shallow copy and will **not** copy object elements within arrays. This is mainly important for arrays containing Python objects.

## ▼ Transpose

- To transpose a 2D array (matrix), i.e., **swap rows and columns**, simply use the `T` attribute of an array object. Note that `np.transpose()` and `<ndarray>.transpose()` also yields the same result.

```
import numpy as np

x = np.array([[1, 2], [3, 4]])
print(x)           # Prints [[1 2]
                    #         [3 4]]
print(x.T)         # Prints [[1 3]
                    #         [2 4]]
print(np.transpose(x)) # Prints [[1 3]
                    #         [2 4]]

# Note that taking the transpose of a rank 1 array does nothing:
v = np.array([1, 2, 3])
print(v)           # Prints [1 2 3]
print(v.T)         # Prints [1 2 3]
```

- With `np.transpose()`, you can not only transpose a matrix but also **rearrange the axes** of a multidimensional array in any order.

```
import numpy as np

a = np.ones((1, 2, 3))
```

```
print(a.shape) # Prints (1, 2, 3)
print(np.transpose(a, (1, 0, 2)).shape) # Prints (2, 1, 3)
```

## ▼ Element-wise sign

- Returns an element-wise sign of an array's elements.

```
import numpy as np

a = np.array([[1, 2], [-3, -4]])

print(np.sign(a)) # Prints [[ 1  1]
                  #         [-1 -1]]
```

## ▼ Sum

- Sum of an array's elements.

```
import numpy as np

a = np.array([[1, 2], [3, 4]])

# Compute sum of all elements
print(np.sum(a)) # Prints 10

# Compute sum of each column
print(np.sum(a, axis=0)) # Prints [4 6]

# Compute sum of each row
print(np.sum(a, axis=1)) # Prints [3 7]
```

## ▼ Average/Mean

- Recall that the mean is the sum of the elements divided by the length of the vector.

```
import numpy as np

a = np.array([[1, -1], [2, -2], [3, -3]]) # Define a 3x2 matrix. Chosen to be a matrix with 0 mean.

# Get the mean for the whole matrix
print(np.mean(a)) # Prints 0.0

# Get the mean for each column. Returns 2 elements.
print(a, axis=0) # Prints [ 2. -2.]

# Get the mean for each row. Returns 3 elements.
print(a, axis=1) # Prints [0. 0. 0.]
```

## ▼ Product

- Return the product of array elements over a given axis.

```
import numpy as np

a = np.array([[1, 2], [3, 4]])

print(np.prod(a)) # Prints 24
```

## ▼ Max

- Return the maximum element of an array or maximum along an axis.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])

print(np.max(a)) # Prints 6
```

## ▼ Element-wise max

- Compare two compatible arrays and return their element-wise maximum. Here, 'compatible' means that one array can be broadcast to the other.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])
b = np.array([[7, 8, 9], [-2, -1, 0]])

print(np.maximum(a, b)) # Prints [[7 8 9]
                        #          [4 5 6]]
```

## ▼ Argmax

- Return the indices of the maximum values along an axis.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])

print(np.argmax(a))          # Prints 5
print(np.argmax(a, axis=0)) # Prints [1 1 1]
print(np.argmax(a, axis=1)) # Prints [2 2]
```

- Note that in case of multiple occurrences of the maximum values, only the index corresponding to the **first occurrence** is returned.

## ▼ Argmin

- Return the indices of the minimum values along an axis.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])

print(np.argmin(a))          # Prints 0
print(np.argmin(a, axis=0))  # Prints [0 0 0]
print(np.argmin(a, axis=1))  # Prints [0 0]
```

- Note that in case of multiple occurrences of the minimum values, only the index corresponding to the **first occurrence** is returned.

## ▼ Where

- Return elements chosen from *x* or *y* depending on condition.
- `np.where()` accepts three arguments: (*condition*, *x*, *y*). It first evaluates the *condition* argument – if it returns `True`, it yields *x*, otherwise yields *y*.

```
import numpy as np

a = np.array([0, 1, 2, 3, 4, 5])

print(np.where(a < 3, a, 10*a)) # Prints [ 0  1  2 30 40 50]
```

## ▼ Pad

- Pad an array.
- `np.pad()` accepts three arguments: *array*, *pad\_width*, *mode='constant'*.
  - *pad\_width* specifies the number of values padded to the edges of each axis. It accepts both an `int` and a `tuple`.
    - `(pad,)` or simply `a pad` (as an `int`) is a shortcut for `before = after = pad width` for all axes.
    - `((before_1, after_1), ... (before_N, after_N))` specifies the padding widths for each axis.
    - `((before, after),)` yields same before and after padding width for each axis.

- When mode is set to constant (which is default), an optional parameter constant\_values lets you specify the padding values for each axis, which has a default of 0 (i.e., zero padding). Rather than just an int, constant\_values accepts a tuple as well:
  - (constant,) or simply a constant (as an int) is a shortcut for before = after = constant for all axes. Default is 0.
  - ((before\_1, after\_1), ... (before\_N, after\_N)) specifies the unique pad constants for each axis.
  - ((before, after),) yields same before and after constants for each axis.

```
import numpy as np
```

```
a = np.array([1, 2, 3, 4, 5])
print(np.pad(a, 1, 'constant'))           # Prints [0 1 2 3 4 5 0]
print(np.pad(a, (2, 3), 'constant', constant_values=(4, 6))) # Prints [4 4 0 1 2 3 4 5 6 6 6]
```

```
a = np.array([[1, 2], [3, 4]])
print(np.pad(a, 1, 'constant', constant_values=(4, 6))) # Prints [[4 4 4 6]
#                                                         [4 1 2 6]
#                                                         [4 3 4 6]
#                                                         [4 6 6 6]]
```

## ▼ Ravel

- Return a **contiguous** flattened array.
- Note that np.ravel() returns a **view** of the array when the elements are contiguous in memory (just like [np.reshape\(.\)](#)). So modifying the result of np.ravel() would also modify the original array. However, np.ravel() returns a copy if, for e.g., the input array were made from slicing another array using a **non-unit** step size (e.g. a = x[:, :2]).

```
import numpy as np
```

```
a = np.array([[1, 2, 3], [4, 5, 6]])
print(np.ravel(a)) # Prints [1, 2, 3, 4, 5, 6]
```

- Note that `np.ravel()` is equivalent to `np.reshape(-1)`.

```
import numpy as np
```

```
a = np.array([[1, 2, 3], [4, 5, 6]])
print(np.reshape(a, -1)) # Prints [1, 2, 3, 4, 5, 6]
```

## ▼ Unravel Index

- Converts a **"flat" index** (i.e., an index into the flattened version of an array) into a **tuple of coordinate arrays**.
- Background:
  - Computer memory is addressed linearly. Each memory cell corresponds to a number. A block of memory can be addressed in terms of a base, which is the memory address of its first element, and the item index. For example, assuming the base address is 10,000:

```
...
item index      0      1      2      3
memory address 10,000 10,001 10,002 10,003
...
```

- To store multi-dimensional blocks, their geometry must somehow be made to fit into linear memory. In C and NumPy, this is done row-

```
...
  | 0      1      2      3
--+-----+
0 | 0      1      2      3
1 | 4      5      6      7
2 | 8      9     10     11
...
```

- So, for example, in this 3-by-4 block the 2D index `((1, 2))` would correspond to the linear index `(6)` which is `(1 * 4 + 2)`.
- `np.unravel_index()` does the inverse. Given a linear index, it computes the corresponding coordinates. Since this depends on the shape of the array, it is a function of the array shape.

```
import numpy as np

# Passing an integer array whose elements are indices into the flattened version of an array
print(np.unravel_index([1, 2, 3], (4, 5)) # Prints (array([0, 0, 0]), array([1, 2, 3]))

# You're not limited to the 2-dimensional XY co-ordinate space
print(np.unravel_index(7, (1, 1, 7)))      # Prints (0, 0, 6)
```

## ▼ Exponential

- Calculate the exponential of all elements in the input array.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])

print(np.exp(a)) # Prints [[ 2.71828183  7.3890561 20.08553692]
                    #       [ 54.59815003 148.4131591 403.42879349]]
```

## ▼ Unique

- Find the unique elements of an array.

```
import numpy as np

a = np.array([[1, 1, 1], [2, 2, 2]])

print(np.unique(a)) # Prints [1 2]
```

## ▼ Bincount

- Count number of occurrences of each value in a one-dimensional array of **non-negative integers**.



```
import numpy as np

a = np.array([1, 1, 1, 2, 2, 2])

print(np.bincount(a)) # Prints [0 3 3]
```

## ▼ Element-wise square

- Return the element-wise square of the input.

```
import numpy as np

a = np.array([[1, 2, 3], [4, 5, 6]])

print(np.square(a)) # Prints [[ 1  4  9]
                      #       [[16 25 36]]
```

## ▼ Element-wise square root

- Return the element-wise non-negative square root of an array.

```
import numpy as np

a = np.array([[1, 4, 9], [16, 25, 36]])

print(np.sqrt(a)) # Prints [[1. 2. 3.]
                      #       [4. 5. 6.]]
```

## ▼ Split

- Split an array into multiple sub-arrays.

```
import numpy as np

a = np.arange(8)

print(np.split(a, 2)) # Prints [array([0, 1, 2, 3]), array([4, 5, 6, 7])]
```

## ▼ Horizontal split

- Split an array into multiple sub-arrays horizontally (column-wise).

```
import numpy as np

a = np.arange(4).reshape(2, 2)

print(a)           # Prints [[0 1]
                    #          [2 3]]

print(np.hsplit(a, 2)) # Prints [[0],
                                #          [2]],
                                #          [[1],
                                #          [3]]
```

## ▼ Vertical split

- Split an array into multiple sub-arrays vertically (row-wise).

```
import numpy as np

a = np.arange(4).reshape(2, 2)

print(a)           # Prints [[0 1]
                    #          [2 3]]
```

```
print(np.vsplit(a, 2)) # Prints [array([[0, 1]]), array([[2, 3]])]
```

## ▼ Stack

- Join a sequence of arrays along a new axis. Note that by default, `axis` is set to 0 implying it joins arrays row-wise.

```
import numpy as np
```

```
a = np.array([1, 2, 3])
```

```
b = np.array([4, 5, 6])
```

```
print ((a,b))                # Prints (array([1, 2, 3]), array([4, 5, 6]))
```

```
print(np.stack((a, b)))      # Prints [[1 2 3],  
                             #         [4 5 6]]
```

```
# This is the same as above since axis is set to 0 by default
```

```
print(np.stack((a, b), axis=0)) # Prints [[1 2 3],  
                                         #         [4 5 6]]
```

```
# np.vstack() matches the output of np.stack().
```

```
# Note that np.vstack() is discussed in more detail below.
```

```
print(np.vstack(a, b))      # Prints [[1 2 3],  
                             #         [4 5 6]]
```

```
# Joining arrays column-wise.
```

```
print(np.stack((a, b), axis=1)) # Prints [[1 4],  
                                         #         [2 5],  
                                         #         [3 6]]
```

```
# np.hstack() does *not* match the output of np.stack().
```

```
print(np.hstack((a, b)))    # Prints [1 2 3 4 5 6]
```

## ▼ Horizontal stack

- Stack arrays in sequence horizontally (column wise).
- Note that `np.hstack()` does not accept an `axis` argument.

```
import numpy as np

a = np.array([1, 2, 3])
b = np.array([4, 5, 6])

print(np.hstack((a, b))) # Prints [1 2 3 4 5 6]
```

## ▼ Vertical stack

- Stack arrays in sequence vertically (row wise).
- Note that `np.vstack()` does not accept an `axis` argument.

```
import numpy as np

a = np.array([1, 2, 3])
b = np.array([4, 5, 6])

print(np.vstack((a, b))) # Prints [[1 2 3],
                          #         [4 5 6]]
```

## ▼ Concatenate

- Join a sequence of arrays along an existing axis.

```
import numpy as np
```

```
a = np.array([[1, 2], [3, 4]])
b = np.array([[5, 6]])

print(np.concatenate((a, b), axis=0)) # Prints [[1 2],
                                         #          [3 4],
                                         #          [5 6]])
```

## ▼ Test any

- Test whether any array element along a given axis evaluates to `True`.
- Returns a single boolean unless `axis` is not `None` (which is the default).
- Note that not a number (NaN),  $\infty$  and  $-\infty$  all evaluate to `True` because these are not equal to 0.

```
import numpy as np

a = [[True, False], [True, True]]
print(np.any(a)) # Prints True

a = [[True, False], [True, True]]
print(np.any(a, axis=0)) # Prints [ True,  True]

a = np.array([[1, 2], [3, 4]])
print(np.any(a)) # Prints True

# A nifty use-case of np.any() is to select elements that satisfy
# atleast one condition from multiple given conditions
a = np.array([[1, 2], [3, 4]])
print(np.any([a > 1, a < 5], axis=0)) # Prints [[ True  True]
                                         #          [ True  True]]
print(a[np.any([a > 1, a < 5], axis=0)]) # Prints [1 2 3 4]

# If axis is None (default), np.any() yields a single boolean by performing
# a logical OR reduction over all the dimensions of the input array
print(np.any([a > 1, a < 5])) # Prints True
```

## ▼ Test all

- Test whether all array elements along a given axis evaluate to `True`.
- Similar to `np.any()`, `np.all()` returns a single boolean unless `axis` is not `None` (which is the default).
- Again, similar to `np.any()`, note that not a number (NaN),  $\infty$  and  $-\infty$  all evaluate to `True` because these are not equal to 0.

```
import numpy as np

a = [[True, False], [True, True]]
print(np.all(a))                # Prints True

a = [[True, False], [True, True]]
print(np.all(a, axis=0))        # Prints [ True, False]

a = np.array([[1, 2], [3, 4]])
print(np.all(a))                # Prints True

# A nifty use-case of np.all() is to select elements that satisfy
# all given conditions
a = np.array([[1, 2], [3, 4]])
print(np.all([a > 1, a < 5], axis=0))  # Prints [[False  True]
                                     #           [ True  True]]
print(a[np.all([a > 1, a < 5], axis=0)]) # Prints [2 3 4]

# Note that if only one condition is specified, np.any() and np.all()
# are equivalent, and have identical outputs
print(a[np.any([a > 1], axis=0)])      # Prints [2 3 4]
print(a[np.all([a > 1], axis=0)])      # Prints [2 3 4]

# If axis is None (default), np.all() yields a single boolean by performing
# a logical AND reduction over all the dimensions of the input array
print(np.all([a > 1, a < 5]))          # Prints False
```

## References and credits

- Parts of this tutorial were originally contributed by [Justin Johnson](#).
- [Stanford CS231n Python/NumPy Tutorial](#) served as a major inspiration for this tutorial.
- [NumPy documentation: Broadcasting](#)
- [What is an intuitive explanation of np.unravel\\_index?](#)
- [When should I use hstack/vstack vs. append vs. concatenate vs. column\\_stack?](#)
- [What is the difference between NumPy's array\(\) and asarray\(\) functions?](#)
- [What does -1 mean in NumPy reshape?](#)
- [Find the most frequent number in a NumPy vector](#)
- [NumPy: From ND to 1D arrays](#)
- [Difference between functions generating random numbers in NumPy](#)
- [How do I select elements of an array given condition?](#)

## Citation

If you found our work useful, please cite it as:

```
@article{Chadha2020DistilledNumPyTutorial,  
  title   = {NumPy Tutorial},  
  author  = {Chadha, Aman},  
  journal = {Distilled AI},  
  year    = {2020},  
  note    = {\url{https://aman.ai}}  
}
```

[Colab paid products](#) - [Cancel contracts here](#)

